

Distributional Reinforcement Learning with Function Approximation

Abstract

These notes give an overview of some theoretical foundations behind approximation based distributional reinforcement learning and newer methods using differentiable metrics on probability measures. We begin with linear value-function approximation in the standard scalar case, then we replace scalar value functions by the entire return distribution and introduce the two classic representations used in this case, quantile and categorical, of probability distributions. Then, we go over how these representations can be used in function approximation using deep neural networks. Finally, we cover the recent research on distributional reinforcement learning using differentiable metrics on probability measures such as the MMD distances or Sinkhorn divergence.

Preliminaries

The first part of these notes follow the treatment of distributional reinforcement learning in [BDR23]. We first recall some of the preliminary definitions and results.

Definition 1: Finite Markov decision process and policy

A finite Markov decision process is given by

$$(\mathcal{X}, \mathcal{A}, \xi_0, P_X, P_R),$$

where $\mathcal{X}$ is a finite state space, $\mathcal{A}$ is a finite action space, ξ_0 is an initial-state distribution, $P_X(\cdot \mid x, a)$ is the next-state kernel, and $P_R(\cdot \mid x, a)$ is the reward distribution. A policy is a map

$$\pi : \mathcal{X} \rightarrow \mathcal{P}(\mathcal{A}), \quad x \mapsto \pi(\cdot \mid x),$$

and the discount factor satisfies $\gamma \in [0, 1)$. The markov chain induced by P_X and π gives rises to the trajectory measure P_π .

Given a policy π , the trajectory process is

$$\begin{aligned} X_0 &\sim \xi_0, & A_t &\sim \pi(\cdot \mid X_t), \\ R_t &\sim P_R(\cdot \mid X_t, A_t), & X_{t+1} &\sim P_X(\cdot \mid X_t, A_t), & t \geq 0. \end{aligned}$$

Starting from $X_0 = x$ and then following π , define the discounted return

$$G^\pi(x) = \sum_{t=0}^{\infty} \gamma^t R_t.$$

using this return distribution, one can naturally define the value function that tells us for each state, what the expected return is.

Definition 2: Value function and return distribution

The state-value function is the expected return

$$V^\pi(x) = \mathbb{E}_\pi[G^\pi(x)].$$

The return distribution at state x is

$$\eta_\pi(x) = \text{Law}(G^\pi(x)).$$

By definition, we can also write the value function as an integral using the return distribution as a measure.

$$V^\pi(x) = \int_{\mathbb{R}} z \eta_\pi(x)(dz).$$

Central to Reinforcement learning is the Bellman operator and its fixed point equation.

Definition 3: Bellman operator for scalar values

For a function $V : \mathcal{X} \rightarrow \mathbb{R}$, the Bellman operator for policy π is

$$(\mathcal{T}^\pi V)(x) = \mathbb{E}_{a \sim \pi(\cdot | x)} [\mathbb{E}[R + \gamma V(X') \mid X = x, a]].$$

Equivalently, the outer expectation samples the action from the policy and the inner expectation samples reward and next state from the MDP. The Bellman equation is

$$V^\pi = \mathcal{T}^\pi V^\pi.$$

There exists a stronger version of this Bellman equation, which is its distributional variant.

$$A \sim \pi(\cdot \mid x), \quad R \sim P_R(\cdot \mid x, A), \quad X' \sim P_X(\cdot \mid x, A),$$

the return satisfies

$$G^\pi(x) \stackrel{D}{=} R + \gamma G^\pi(X') \quad \text{given } X = x.$$

Taking laws gives

$$\eta_\pi(x) = \text{Law}(R + \gamma G^\pi(X') \mid X = x).$$

Definition 4: Distributional Bellman operator

For $r, z \in \mathbb{R}$, define what in [BDR23] is called the bootstrap map via

$$b_{r,\gamma}(z) = r + \gamma z.$$

If $\nu \in \mathcal{P}(\mathbb{R})$ and $Z \sim \nu$, then the pushforward $(b_{r,\gamma})_{\#}\nu$ is the law of $r + \gamma Z$:

$$(b_{r,\gamma})_{\#}\nu = \text{Law}(r + \gamma Z).$$

For a return-distribution function $\eta : \mathcal{X} \rightarrow \mathcal{P}(\mathbb{R})$, the distributional Bellman operator is

$$(\mathcal{T}^{\pi}\eta)(x) = \mathbb{E}_{a \sim \pi(\cdot|x)} [\mathbb{E} [(b_{R,\gamma})_{\#}\eta(X') \mid X = x, a]].$$

The return-distribution function η_{π} is its fixed point:

$$\eta_{\pi} = \mathcal{T}^{\pi}\eta_{\pi}.$$

For the theory of reinforcement learning particularly important, is that the Bellman operator is a contraction. This is again true for the distributional Bellman operator. However in this case, it is a contraction with respect to the Wasserstein distance.

Theorem 1: Wasserstein contraction of the distributional Bellman operator

For $p \in [1, \infty]$, let

$$\overline{W}_p^{\infty}(\eta, \nu) = \sup_{x,a} W_p(\eta(x, a), \nu(x, a)).$$

Then

$$\overline{W}_p^{\infty}(\mathcal{T}^{\pi}\eta, \mathcal{T}^{\pi}\nu) \leq \gamma \overline{W}_p^{\infty}(\eta, \nu).$$

As also covered earlier in the book, arbitrary distributions cannot be stored on the computer. Hence we need to find some finite dimensional parametrized approximation. The two main representations are the following.

Definition 5: Finite representations of return distributions

For a probability measure $\eta \in \mathcal{P}(\mathbb{R})$, let

$$F_{\eta}(t) = \eta((-\infty, t]), \quad F_{\eta}^{-1}(\tau) = \inf\{z \in \mathbb{R} : F_{\eta}(z) \geq \tau\}.$$

The two standard representations are:

$$\text{categorical:} \quad \eta(x) = \sum_{i=1}^m p_i(x) \delta_{\theta_i}, \quad p_i(x) \geq 0, \quad \sum_{i=1}^m p_i(x) = 1,$$

and

$$\text{quantile:} \quad \eta(x) = \frac{1}{m} \sum_{i=1}^m \delta_{\theta_i(x)}, \quad \tau_i = \frac{2i-1}{2m}, \quad \theta_i(x) \approx F_{\eta(x)}^{-1}(\tau_i).$$

If $\mathcal{X}$ is large, storing a scalar value for every state is already expensive. In our case, we would need to store a full probability distribution for every state which is even more expensive. Hence, the rest of the notes which are based on [BDR23] will focus on value and action value function approximation, starting with the linear case.

1 Linear Function Approximation

We first discuss linear approximation for the state-value function

$$V^\pi : \mathcal{X} \rightarrow \mathbb{R}.$$

This section follows parts of Chapter 9 of [BDR23].

1.1 Function Approximation

In linear function approximation, we represent the value function as a linear combination of features. Similar to neural networks, the features here are allowed to be nonlinear.

Definition 6: State representation and linear value approximation

Let $n \in \mathbb{N}$. A state representation is a mapping

$$\phi : \mathcal{X} \rightarrow \mathbb{R}^n, \quad \phi(x) = \begin{pmatrix} \phi_1(x) \\ \vdots \\ \phi_n(x) \end{pmatrix}.$$

For $w \in \mathbb{R}^n$, the linear value function approximation is

$$V_w(x) = \phi(x)^\top w,$$

where each ϕ_i is a basis function and $\phi_i(x)$ is a feature.

The model is linear in its weights and

$$V_{\alpha w_1 + \beta w_2} = \alpha V_{w_1} + \beta V_{w_2}, \quad \nabla_w V_w(x) = \phi(x).$$

For finite $\mathcal{X}$, collect the row vectors $\phi(x)^\top$ in a feature matrix $\Phi \in \mathbb{R}^{|\mathcal{X}| \times n}$. Then

$$V_w = \Phi w, \quad \mathcal{F}_\phi = \{\Phi w : w \in \mathbb{R}^n\}.$$

The same can be done for action value function using $\phi : \mathcal{X} \times \mathcal{A} \rightarrow \mathbb{R}^n$ and $Q_w(x, a) = \phi(x, a)^\top w$. Since we want to focus on what the distributional part changes about the approximation later, we will mostly discuss the state value case.

1.2 Optimal Linear Value Function Approximation

As states above, the equation

$$V_w = \Phi w$$

defines a linear system, which won't have a solution in general. Instead, a standard way to find the 'best' solution is to minimize the squared error. This can also be done here, however in [BDR23] it is mentioned that in RL we are often not equally interested in all states, hence we introduce a weighted least squares loss.

So, let $\xi \in \mathcal{P}(\mathcal{X})$ be the relative importance of the states and define

$$\|V - V'\|_{\xi,2} = \left(\sum_{x \in \mathcal{X}} \xi(x) (V(x) - V'(x))^2 \right)^{1/2}.$$

we are thus interested in the best approximation to V^π which solves

$$\min_{w \in \mathbb{R}^n} \|V^\pi - V_w\|_{\xi,2}.$$

and this is given by a weighted normal equation.

Theorem 2: Weighted normal equation

Let $\Xi = \text{diag}(\xi(x) : x \in \mathcal{X})$. If ξ has full support and the columns of Φ are linearly independent, the unique solution is

$$w^* = (\Phi^\top \Xi \Phi)^{-1} \Phi^\top \Xi V^\pi.$$

Thus the state weighting ξ is part of the approximation problem: changing how often states are sampled generally changes the best attainable weights.

The problem is, we can't iterate the Bellman operator since after applying it, we might not have a linear representation anymore. This motivates the following definition of a projected Bellman operator.

1.3 The Projected Bellman Operator

Definition 7: Projected Bellman operator

Define the weighted orthogonal projection $\Pi_{\phi,\xi}$ by

$$(\Pi_{\phi,\xi} V)(x) = \phi(x)^\top w^*, \quad w^* \in \arg \min_{w \in \mathbb{R}^n} \|V - V_w\|_{\xi,2}.$$

Approximate dynamic programming iterates

$$V_{k+1} = \Pi_{\phi,\xi} \mathcal{T}^\pi V_k.$$

Luckily, one can still show that this projected Bellman operator is a contraction and that the fixed point of the projected Bellman operator is close to the best approximation of the

true value function. Hence, by iteration the Bellman equation like this, we will converge to something that is close to what we are trying to approximate!

Theorem 3: Projected Bellman contraction and error bound

Let ξ_π be a fully supported steady-state distribution under π . Then Π_{ϕ, ξ_π} is a nonexpansion in the ξ_π -weighted L_2 norm. Consequently,

$$\|\Pi_{\phi, \xi_\pi} \mathcal{T}^\pi V - \Pi_{\phi, \xi_\pi} \mathcal{T}^\pi V'\|_{\xi_\pi, 2} \leq \gamma \|V - V'\|_{\xi_\pi, 2}.$$

The projected operator has a unique fixed point

$$\hat{V}^\pi = \Pi_{\phi, \xi_\pi} \mathcal{T}^\pi \hat{V}^\pi,$$

Additionally, they show that it holds

$$\|\hat{V}^\pi - V^\pi\|_{\xi_\pi, 2} \leq \frac{1}{\sqrt{1 - \gamma^2}} \|\Pi_{\phi, \xi_\pi} V^\pi - V^\pi\|_{\xi_\pi, 2}.$$

which means that the fixed point of the projected Bellman operator is close to the best approximation.

Proof. The proof follows the argument from [BDR23, Theorem 9.8]. The transition operator P^π is nonexpansive and due to the orthogonality of the projection we have for every U ,

$$U = \Pi_{\phi, \xi_\pi} U + (I - \Pi_{\phi, \xi_\pi})U$$

is an orthogonal decomposition in the ξ_π -weighted inner product. Hence, by Pythagoras,

$$\|U\|_{\xi_\pi, 2}^2 = \|\Pi_{\phi, \xi_\pi} U\|_{\xi_\pi, 2}^2 + \|(I - \Pi_{\phi, \xi_\pi})U\|_{\xi_\pi, 2}^2,$$

so $\|\Pi_{\phi, \xi_\pi} U\|_{\xi_\pi, 2} \leq \|U\|_{\xi_\pi, 2}$. Thus the composition with $\mathcal{T}^\pi$ is a γ -contraction. Banach's fixed-point theorem gives a unique fixed point and convergence of iterating the projected Bellman operator.

For the error estimate, we use orthogonality of the projection and $V^\pi = \mathcal{T}^\pi V^\pi$ which yields

$$\begin{aligned} \|\hat{V}^\pi - V^\pi\|_{\xi_\pi, 2}^2 &= \|\hat{V}^\pi - \Pi_{\phi, \xi_\pi} V^\pi\|_{\xi_\pi, 2}^2 + \|\Pi_{\phi, \xi_\pi} V^\pi - V^\pi\|_{\xi_\pi, 2}^2 \\ &\leq \gamma^2 \|\hat{V}^\pi - V^\pi\|_{\xi_\pi, 2}^2 + \|\Pi_{\phi, \xi_\pi} V^\pi - V^\pi\|_{\xi_\pi, 2}^2. \end{aligned}$$

Rearranging and taking square roots proves the bound. $\square$

However, while this result is nice, performing full Bellman update is expensive. A more practical tool is TD learning which can also be done in the approximate regime. Since we again want to focus on the distributional side, we don't mention it here and refer to [BDR23, Chapter 9.4] for more details.

2 Distributional Value Approximation

Distributional reinforcement learning replaces the scalar value function by a return-distribution function. In the following, we use the action-value version of the distributional Bellman op-

erator defined in the preliminaries.

As highlighted in the prelims, the book mainly discusses two ways to parametrize the return distribution. In both cases we are talking about probability distribution, which do not form a linear space. This is a problem, since we cannot easily define a orthogonal projection. Hence, for now, we focus on the TD learning of the distributional case. First, we define the linear approximation for both cases.

Definition 8: Categorical return-distribution approximation

Fix the support of the atoms

$$S = \{\theta_1 < \dots < \theta_m\} \subset \mathbb{R}.$$

The categorical approximation keeps these locations fixed and learns their probabilities:


$$\eta_W(x, a) = \sum_{i=1}^m p_i(x, a; W) \delta_{\theta_i}.$$

to learn probabilities, one classically uses the softmax function of the approximated function. Hence, we have $\ell_i(x, a) = \phi(x, a)^\top w_i$, as our parametrization and

$$p_i(x, a; W) = \frac{\exp(\phi(x, a)^\top w_i)}{\sum_{j=1}^m \exp(\phi(x, a)^\top w_j)}.$$

Categorical: fixed locations, learned masses

$$\eta(x) = \sum_{i=1}^m p_i(x) \delta_{\theta_i}$$



Quantile: equal masses, learned locations

$$\eta(x) = \frac{1}{m} \sum_{i=1}^m \delta_{\theta_i(x)}$$




Figure 1: Comparison between categorical and quantile approximations.

With the categorical approximation, we have the problem that after a Bellman or even TD update, we might not live on S anymore. Hence we need the following projection.

Definition 9: Categorical projection Π_c

The Bellman update moves the fixed atoms because

$$(b_{r,\gamma})_{\#}\eta(x') = \sum_{j=1}^m p_j(x')\delta_{r+\gamma\theta_j}, \quad b_{r,\gamma}(z) = r + \gamma z.$$

In general, the shifted locations $r + \gamma\theta_j$ are not elements of the fixed support S . The categorical projection $\Pi_c : \mathcal{P}(\mathbb{R}) \rightarrow \mathcal{P}(S)$ is defined first on a Dirac mass. If $\theta_\ell \leq y \leq \theta_{\ell+1}$, then

$$\Pi_c \delta_y = \frac{\theta_{\ell+1} - y}{\theta_{\ell+1} - \theta_\ell} \delta_{\theta_\ell} + \frac{y - \theta_\ell}{\theta_{\ell+1} - \theta_\ell} \delta_{\theta_{\ell+1}}.$$

Values outside the support are clipped to its endpoints:

$$\Pi_c \delta_y = \begin{cases} \delta_{\theta_1}, & y \leq \theta_1, \\ \delta_{\theta_m}, & y \geq \theta_m. \end{cases}$$

The definition is extended linearly to discrete measures. Thus, if $\theta_\ell \leq r + \gamma\theta_j \leq \theta_{\ell+1}$, the mass $p_j(x')$ contributes

$$\begin{aligned} \bar{p}_\ell &+= p_j(x') \frac{\theta_{\ell+1} - (r + \gamma\theta_j)}{\theta_{\ell+1} - \theta_\ell}, \\ \bar{p}_{\ell+1} &+= p_j(x') \frac{(r + \gamma\theta_j) - \theta_\ell}{\theta_{\ell+1} - \theta_\ell}. \end{aligned}$$

The projected categorical Bellman target is therefore

$$\Pi_c(b_{r,\gamma})_{\#}\eta(x') = \sum_{i=1}^m \bar{p}_i \delta_{\theta_i}.$$

The quantile family instead parametrizes $\nu \in \mathcal{P}(\mathbb{R})$ with cumulative distribution function F_ν , through its quantile function by

$$F_\nu^{-1}(\tau) = \inf\{z \in \mathbb{R} : F_\nu(z) \geq \tau\}, \quad \tau \in (0, 1).$$

in the quantile representation, we try and approximate the positions of fixed quantile levels.

Definition 10: Quantile return-distribution approximation

The general form of linear quantile return distribution approximation is

$$\theta_i(x, a) = \phi(x, a)^\top w_i, \quad w_i \in \mathbb{R}^d.$$

Equivalently, with $W = (w_1, \dots, w_N) \in \mathbb{R}^{d \times N}$,

$$\eta_W(x, a) = \frac{1}{N} \sum_{i=1}^N \delta_{\phi(x, a)^\top w_i}.$$

In particular, we fix the probabilities to be equal, but allow the atom positions to change.

2.1 Signed categorical approximation

For the theoretical analysis of linear categorical approximation, we now want to allow signed measures. This will allow us to analyze the space of measures as a linear space equipped with a distance and as such also will allow us to define a linear projection operator. In turn this will in the end allow us to see if in the linear categorical setting, we are working with a contraction (if we work in the categorical representation defined below).

Definition 11: Signed categorical family

For the fixed support $S = \{\theta_1, \dots, \theta_m\}$, define

$$\mathcal{F}_{S,m} = \left\{ \sum_{i=1}^m p_i \delta_{\theta_i} \mid p_i \in \mathbb{R}, \sum_{i=1}^m p_i = 1 \right\}.$$

The coefficients still have total mass one but may be negative. The associated return-distribution functions form the space of maps from the state space to these signed measures:

$$\mathcal{F}_{S,m}^{\mathcal{X}} = \{\eta : \mathcal{X} \rightarrow \mathcal{F}_{S,m}\}.$$

Now, we will not use the softmax anymore, but instead we will just directly use $\phi(x)^\top w_i$ as the probabilities, although centering them to ensure that they sum to one.

$$p_i(x) = \frac{1}{m} + \phi(x)^\top w_i - \frac{1}{m} \sum_{j=1}^m \phi(x)^\top w_j.$$

Then $\sum_i p_i(x) = 1$ for every x , while individual $p_i(x)$ may be negative. We denote the resulting finite dimensional measure space by

$$\mathcal{F}_{\phi,S,m} = \left\{ x \mapsto \sum_{i=1}^m p_i(x) \delta_{\theta_i} \mid w_1, \dots, w_m \in \mathbb{R}^d \right\}.$$

Now, we need a distance on this space, which the Cramer distance defined earlier in the book provides.

Definition 12: Cramér geometry for signed measures

For signed measures ν, ν' of total mass one, define

$$\ell_2^2(\nu, \nu') = \int_{\mathbb{R}} (F_\nu(z) - F_{\nu'}(z))^2 dz, \quad F_\nu(z) = \nu((-\infty, z]).$$

Let

$$\mathcal{M}_{\ell_2,1}(\mathbb{R}) = \{\nu : \nu(\mathbb{R}) = 1, \ell_2(\nu, \delta_0) < \infty\}.$$

For return-distribution functions, the stationary-state-weighted distance is

$$\ell_{\xi,2}^2(\eta, \eta') = \sum_{x \in \mathcal{X}} \xi(x) \ell_2^2(\eta(x), \eta'(x)).$$

Now, there are two distinct approximation errors after a Bellman update. First, $(\mathcal{T}^\pi \eta)(x)$ need not be supported on S , so the categorical projection Π_c moves its mass back to the fixed support. Second, even $\Pi_c \mathcal{T}^\pi \eta$ need not have probabilities that are linear in the features, so a feature projection is required:

$$\Pi_{\phi, \xi, \ell_2} \tilde{\eta} \in \arg \min_{\eta' \in \mathcal{F}_{\phi, S, m}} \ell_{\xi,2}(\tilde{\eta}, \eta').$$

Consequently, the approximate Bellman operator we need to consider here is the double projection

$$\Pi_{\phi, \xi, \ell_2} \Pi_c \mathcal{T}^\pi.$$

Lemma 1: Bellman and categorical bounds [BDR23, Lemma 9.14]

Let ξ_π be a fully supported stationary distribution under π . For $\eta, \eta' \in \mathcal{M}_{\ell_2,1}(\mathbb{R})^{\mathcal{X}}$,

$$\ell_{\xi_\pi,2}(\mathcal{T}^\pi \eta, \mathcal{T}^\pi \eta') \leq \sqrt{\gamma} \ell_{\xi_\pi,2}(\eta, \eta').$$

Moreover, the pointwise categorical projection is nonexpansive:

$$\ell_{\xi_\pi,2}(\Pi_c \eta, \Pi_c \eta') \leq \ell_{\xi_\pi,2}(\eta, \eta').$$

Lemma 2: Linear feature projection nonexpansive [BDR23, Lemma 9.16]

For $\eta, \eta' \in \mathcal{F}_{S,m}^{\mathcal{X}}$,

$$\ell_{\xi_\pi,2}(\Pi_{\phi, \xi_\pi, \ell_2} \eta, \Pi_{\phi, \xi_\pi, \ell_2} \eta') \leq \ell_{\xi_\pi,2}(\eta, \eta').$$

With these two lemmas in hand, the final result will just be a simple combination of the obtained bounds which finally shows that the double projected Bellman operator is a contraction with the same $\sqrt{\gamma}$ constant as the original Bellman operator w.r.t to the cramer geometry.

Theorem 4: Doubly projected categorical Bellman contraction [BDR23, Theorem 9.13]

The operator

$$\Pi_{\phi, \xi_\pi, \ell_2} \Pi_c \mathcal{T}^\pi$$

is a $\sqrt{\gamma}$ -contraction in $\ell_{\xi_\pi, 2}$. Therefore it has a unique fixed point in $\mathcal{F}_{\phi, S, m}$.

Proof. The Bellman operator contributes the strict $\sqrt{\gamma}$ factor, while both projections are nonexpansive. Hence

$$\begin{aligned} & \ell_{\xi_\pi, 2}(\Pi_{\phi, \xi_\pi, \ell_2} \Pi_c \mathcal{T}^\pi \eta, \Pi_{\phi, \xi_\pi, \ell_2} \Pi_c \mathcal{T}^\pi \eta') \\ & \leq \ell_{\xi_\pi, 2}(\Pi_c \mathcal{T}^\pi \eta, \Pi_c \mathcal{T}^\pi \eta') \\ & \leq \ell_{\xi_\pi, 2}(\mathcal{T}^\pi \eta, \mathcal{T}^\pi \eta') \\ & \leq \sqrt{\gamma} \ell_{\xi_\pi, 2}(\eta, \eta'). \end{aligned}$$

Banach's fixed-point theorem gives existence and uniqueness of the fixed point. $\square$

3 Distributional Approximate TD

The categorical TD learning is very similar to the scalar case, hence we focus on the quantile case and refer to [BDR23, Chapter 10].

For the quantile case, if we are given a policy π , sample a transition and next action

$$(x_t, a_t, r_t, x'_t, a'_t), \quad a'_t \sim \pi(\cdot \mid x'_t).$$

The distributional Bellman target is represented by the N target atoms

$$y_j = r_t + \gamma \theta_j(x'_t, a'_t), \quad j = 1, \dots, N.$$

The current prediction at (x_t, a_t) has atoms

$$z_i = \theta_i(x_t, a_t), \quad i = 1, \dots, N.$$

Definition 13: Quantile TD loss

Quantile regression uses the asymmetric absolute loss

$$\rho_\tau(u) = u(\tau - \mathbb{1}_{\{u < 0\}}).$$

The empirical Quantile TD loss is then

$$L(W) = \frac{1}{N} \sum_{i=1}^N \sum_{j=1}^N \rho_{\tau_i}(y_j - z_i).$$

The residual $y_j - z_i$ compares an atom of the Bellman target distribution with the predicted τ_i -quantile at the current state and chosen action. Minimizing this loss moves $\eta_W(x_t, a_t)$ towards the bellman target

$$\frac{1}{N} \sum_{j=1}^N \delta_{r_t + \gamma \theta_j(x'_t, a'_t)} = (b_{r_t, \gamma})_{\#} \eta_W(x'_t, a'_t).$$

For linear quantiles, a subgradient (which we need since its not differentiable at 0) update is

$$w_{i,t+1} = w_{i,t} + \alpha_t \left(\frac{1}{N} \sum_{j=1}^N (\tau_i - \mathbb{1}_{\{y_j < z_i\}}) \right) \phi(x_t, a_t), \quad i = 1, \dots, N.$$

This is the exact analogue of the semi-gradient TD update:

$$\delta_t = r_t + \gamma Q_{w_t}(x'_t, a'_t) - Q_{w_t}(x_t, a_t)$$

If we are not given a policy, we can just greedily sample it at the next state.

$$a^*(x') \in \arg \max_{a' \in \mathcal{A}} \frac{1}{N} \sum_{j=1}^N \theta_j(x', a').$$

The target atoms become

$$y_j = r + \gamma \theta_j(x', a^*(x')), \quad j = 1, \dots, N.$$

4 Deep Distributional Reinforcement Learning

Now, instead of a linear feature representation, we use a neural network

$$\psi_\omega(x) \in \mathbb{R}^d,$$

where ω are the parameters and the final layers produce a distributional prediction in some sense, i.e. either quantiles or probabilities for the categorical case.

4.1 QR-DQN

QR-DQN is the deep version of Quantile TD. It uses a neural network to output N quantile locations for each action:

$$\theta_i(x, a; \omega), \quad i = 1, \dots, N, \quad a \in \mathcal{A}.$$

For each state-action pair, the predicted return distribution is

$$\eta_\omega(x, a) = \frac{1}{N} \sum_{i=1}^N \delta_{\theta_i(x, a; \omega)}.$$

As in DQN, QR-DQN uses two sets of parameters. This is needed because the target network is the same as our current one, hence we fix the target parameters for a couple of iterations always. For a transition (x, a, r, x') , the greedy next action is selected by the mean of the target-network distribution:

$$a^*(x') \in \arg \max_{a' \in \mathcal{A}} \frac{1}{N} \sum_{j=1}^N \theta_j(x', a'; \bar{\omega}).$$

The target quantile atoms are

$$y_j = r + \gamma(1 - d)\theta_j(x', a^*(x'); \bar{\omega}), \quad j = 1, \dots, N.$$

The online prediction atoms for the sampled action are

$$z_i = \theta_i(x, a; \omega), \quad i = 1, \dots, N.$$

Definition 14: Quantile Huber loss

QR-DQN replaces the nonsmooth quantile loss ρ_τ by the quantile Huber loss; see [BDR23, Chapter 10]. Define

$$\mathcal{L}_\kappa(u) = \begin{cases} \frac{1}{2}u^2, & |u| \leq \kappa, \\ \kappa \left(|u| - \frac{1}{2}\kappa \right), & |u| > \kappa, \end{cases}$$

and

$$\rho_\tau^\kappa(u) = \left| \tau - \mathbb{1}_{\{u < 0\}} \right| \frac{\mathcal{L}_\kappa(u)}{\kappa}.$$

For small residuals the loss is quadratic and gives smoother gradients, in particular it is differentiable at 0. For large residuals it behaves linearly, preventing exploding gradients.




Figure 2: Quantile Huber losses for several quantile levels.

Definition 15: QR-DQN training objective

QR-DQN minimizes the pairwise quantile Huber loss

$$L_{\text{QR}}(\omega) = \frac{1}{N} \sum_{i=1}^N \sum_{j=1}^N \rho_{\tau_i}^{\kappa}(y_j - z_i),$$

with quantile levels

$$\tau_i = \frac{2i - 1}{2N}.$$

If we have access to a minibatch $\mathcal{B}$ sampled from replay memory, the objective is just the sum over the above single transition loss.

For the categorical case, the most notable deep variant is C51 which essentially uses the same ideas that the linear approximation in both the scalar and distributional versions use, hence we will not discuss it here. For more details, see [BDR23, Chapter 10.2].

5 Particle Distributional DQN

Now we will discuss some recent work on Distributional DQN that does not utilize a fixed parametrization of the return distribution. Instead, a particle representation is used where we just take the empirical measure of multiple samples.

$$\eta_\omega(x, a) = \frac{1}{N} \sum_{i=1}^N \delta_{z_i(x, a; \omega)}$$

For such methods, the common template is:

1. choose a divergence or metric d between probability measures,
2. study whether the distributional Bellman operator contracts under the induced supremum metric which takes the maximum difference over the states and actions

$$\bar{d}(\eta, \nu) = \sup_{x, a} d(\eta(x, a), \nu(x, a)),$$

3. use the empirical version of d as a differentiable DQN loss.

The last point is important, since the standard Wasserstein distance is not differentiable and hence not that well suited for this problem.

5.1 MMD-DQN

This section follows [NTGV21]. The maximum mean discrepancy (MMD) is a metric on probability measures that is induced by a reproducing kernel Hilbert space. Depending on the choice of the kernel, many different metrics can be constructed. However not all kernels induce a metric.

Definition 16: Kernel mean embedding and maximum mean discrepancy

Let $k : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$ be a continuous positive-definite kernel and let $\mathcal{H}_k$ be its reproducing kernel Hilbert space. For a probability measure μ , define the kernel mean embedding

$$\psi_\mu = \int k(z, \cdot) \mu(dz) \in \mathcal{H}_k.$$

The maximum mean discrepancy is

$$\text{MMD}_k(\mu, \nu) = \|\psi_\mu - \psi_\nu\|_{\mathcal{H}_k}.$$

Equivalently, if $Z, Z' \sim \mu$ and $Y, Y' \sim \nu$ are independent, then

$$\text{MMD}_k^2(\mu, \nu) = \mathbb{E}k(Z, Z') + \mathbb{E}k(Y, Y') - 2\mathbb{E}k(Z, Y).$$

If the kernel is characteristic, the embedding is injective, and $\text{MMD}_k(\mu, \nu) = 0$ implies $\mu = \nu$. With this, the MMD is really a metric.

Example 1: Characteristic kernels

Standard characteristic kernels on $\mathbb{R}^d$ include

$$\text{Gaussian: } k(x, y) = \exp\left(-\frac{\|x - y\|^2}{2\sigma^2}\right),$$

$$\text{Laplace: } k(x, y) = \exp\left(-\frac{\|x - y\|}{\sigma}\right),$$

$$\text{inverse quadratic: } k(x, y) = (c^2 + \|x - y\|^2)^{-\beta},$$

where $\sigma, c, \beta > 0$. By contrast, the linear kernel $k(x, y) = x^\top y$ is not characteristic.

The papers [NTGV21; Sun+24] additionally use

$$k_\alpha(x, y) = -\|x - y\|^\alpha, \quad 0 < \alpha < 2,$$

and call it an *unrectified kernel*. Despite this function not being a kernel, as it is not positive definite it still induces a MMD metric for probability measures. Which we show in the following.

Theorem 5: Unrectified Kernel induces MMD metric

For $0 < \alpha < 2$, the function k_α is conditionally positive definite on coefficients whose sum is zero. Hence, for probability measures μ, ν ,

$$\text{MMD}_{k_\alpha}^2(\mu, \nu) = -\mathbb{E}\|Z - Z'\|^\alpha - \mathbb{E}\|Y - Y'\|^\alpha + 2\mathbb{E}\|Z - Y\|^\alpha \geq 0,$$

where $Z, Z' \sim \mu$ and $Y, Y' \sim \nu$ are independent. In particular, MMD_{k_α} is a metric.

Proof. Most of this proof is utilizing the theory from [Sej+13]. Firstly, one can show directly that

$$\sum_{i,j=1}^N \alpha_i \alpha_j \|x_i - x_j\|^2 \leq 0$$

for any (α_i) with $\sum_{i=1}^N \alpha_i = 0$. This is what [Sej+13] calls a negative type metric. For this, in Proposition 3.1 they show that also $\|\cdot - \cdot\|^\alpha$ is a negative type metric for any $0 < \alpha < 2$. This implies that k_α is conditionally positive definite on coefficients whose sum is zero. Additionally, Lemma 12 of [Sej+13] shows that for any negative-type semimetric, the distance-induced function

$$k_{x_0}(x, y) = \frac{1}{2} (\|x - x_0\|^\alpha + \|y - x_0\|^\alpha - \|x - y\|^\alpha)$$

is a valid kernel and hence induces an RKHS. Additionally, they show that this kernel is characteristic and hence MMD is a metric for this Kernel. When writing out this MMD, we have

$$\text{MMD}_{k_{x_0}}^2(\mu, \nu) = \mathbb{E}k_{x_0}(Z, Z') + \mathbb{E}k_{x_0}(Y, Y') - 2\mathbb{E}k_{x_0}(Z, Y)$$

now when expanding the definition of k_{x_0} , we have

$$\begin{aligned}
MMD_{k_{x_0}}^2(\mu, \nu) &= \frac{1}{2}\mathbb{E}\|Z - x_0\|^\alpha + \frac{1}{2}\mathbb{E}\|Z' - x_0\|^\alpha - \frac{1}{2}\mathbb{E}\|Z - Z'\|^\alpha \\
&\quad + \frac{1}{2}\mathbb{E}\|Y - x_0\|^\alpha + \frac{1}{2}\mathbb{E}\|Y' - x_0\|^\alpha - \frac{1}{2}\mathbb{E}\|Y - Y'\|^\alpha \\
&\quad - \mathbb{E}\|Z - x_0\|^\alpha - \mathbb{E}\|Y - x_0\|^\alpha + \mathbb{E}\|Z - Y\|^\alpha \\
&= \frac{1}{2}(-\mathbb{E}\|Z - Z'\|^\alpha - \mathbb{E}\|Y - Y'\|^\alpha + 2\mathbb{E}\|Z - Y\|^\alpha) \\
&= \frac{1}{2}MMD_{k_\alpha}^2(\mu, \nu)
\end{aligned}$$

The expectations vanish because Z and Z' are independent and identically distributed, hence $\mathbb{E}\|Z - x_0\|^\alpha = \mathbb{E}\|Z' - x_0\|^\alpha$ and the same holds for Y and Y' . The factor two does not affect whether we have a metric or not, hence MMD with the negative type kernel is a metric. $\square$

As we will see below, characteristicness is not required to prove the Bellman contraction as long as we have a kernel like function with two required properties that also makes the MMD a metric. This is why taking a rectified kernel is justified.

Additionally, Nguyen et al. define the state-action supremum metric as hinted at in the introduction [NTGV21]

$$MMD_{\infty, k}(\eta, \nu) = \sup_{x, a} MMD_k(\eta(x, a), \nu(x, a)).$$

Theorem 6: MMD Bellman contraction, [NTGV21, Theorem 2]

Assume that $k = \sum_{i \in I} c_i k_i$ where $c_i \geq 0$ and every k_i has the following two properties. Firstly, it is translation invariant,

$$k(x + c, y + c) = k(x, y)$$

and somewhat scale invariant, i.e. for $b_{\gamma, 0}(z) = \gamma z$,

$$MMD_k((b_{\gamma, 0})_\# \mu, (b_{\gamma, 0})_\# \nu) = \gamma^{\alpha_i/2} MMD_k(\mu, \nu).$$

then, if we set $\alpha_* := \inf_{i \in I} \alpha_i$. The distributional Bellman operator satisfies

$$MMD_{\infty, k}(\mathcal{T}^\pi \eta, \mathcal{T}^\pi \nu) \leq \gamma^{\alpha_*/2} MMD_{\infty, k}(\eta, \nu).$$

In particular, the unrectified distance kernel k_α satisfies the assumptions. On the other hand, the Gaussian kernel does not satisfy the scale invariance since if we look at the scaled gaussian kernel we have

$$k_\sigma(\gamma x, \gamma y) = \exp\left(-\frac{\|\gamma x - \gamma y\|^2}{2\sigma^2}\right) = \exp\left(-\frac{\gamma^2\|x - y\|^2}{2\sigma^2}\right) \neq \gamma^{\alpha/2} k_\sigma(x, y)$$

With the contraction result in hand we can now straight forwardly define the particle DQN method. For a transition (x, a, r, x') and target-network parameters $\bar{\omega}$, the particle DQN

target is constructed by first choosing the greedy next action

$$a^*(x') \in \arg \max_{a' \in \mathcal{A}} \frac{1}{N} \sum_{j=1}^N z_j(x', a'; \bar{\omega})$$

and then shifting the target particles:

$$y_j = r + \gamma z_j(x', a^*(x'); \bar{\omega}), \quad j = 1, \dots, N.$$

With current particles $z_i = z_i(x, a; \omega)$, the empirical MMD loss is evaluated as

$$L_{\text{MMD}}(\omega) = \frac{1}{N^2} \sum_{i, i'=1}^N k(z_i, z_{i'}) + \frac{1}{N^2} \sum_{j, j'=1}^N k(y_j, y_{j'}) - \frac{2}{N^2} \sum_{i, j=1}^N k(z_i, y_j).$$

and this can be differentiated through autograd with respect to the parameters of our networks that produced the samples.

Suprisingly and disagreeing with the theory, [NTGV21] have better empirical performance when we choose the gaussian kernel which means that a distance performs well under which the Bellman Operator is not a contraction.

5.2 SinkhornDRL

As already stated, the Wasserstein distance is not a differentiable metric. This paper [Sun+24] uses a regularized version which makes it work.

Definition 17: Regularized optimal transport and Sinkhorn divergence

For a cost function $c : \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}_+$, the standard OT problem is the following minimization

$$\text{OT}_c(\mu, \nu) = \inf_{\pi \in \Pi(\mu, \nu)} \int c(z, y) \pi(\text{d}z, \text{d}y),$$

where $\Pi(\mu, \nu)$ is the set of couplings with marginals μ and ν . The entropic version is

$$\text{OT}_c^\varepsilon(\mu, \nu) = \inf_{\pi \in \Pi(\mu, \nu)} \left\{ \int c(z, y) \pi(\text{d}z, \text{d}y) + \varepsilon \text{KL}(\pi \mid \mu \otimes \nu) \right\}.$$

However, this version is not a metric anymore since for two measures that are the same, the regularized Wasserstein might be non zero (and it does not fulfill the triangle inequality). This is why the paper uses what they call the debiased Sinkhorn divergence

$$\bar{W}_{c, \varepsilon}(\mu, \nu) = 2 \text{OT}_c^\varepsilon(\mu, \nu) - \text{OT}_c^\varepsilon(\mu, \mu) - \text{OT}_c^\varepsilon(\nu, \nu).$$

While this solves the issue of it being non zero when both entries are the same, this still does not satisfy the triangle inequality.

Based on this $\bar{W}_{c, \varepsilon}$, the paper defines the supremum metric as before

$$\bar{W}_{c, \varepsilon}^\infty(\eta, \nu) = \sup_{x, a} \bar{W}_{c, \varepsilon}(\eta(x, a), \nu(x, a)).$$




Figure 3: Regularization controls how much 'spread out' the mass is, since we bias the optimization towards the product measure which is 'the most spread out'.

Theorem 7: Reg. Wasserstein Bellman contraction, [Sun+24, Theorem 4.1]

Define

$$\bar{\Delta}_\varepsilon(\gamma, \alpha) = \gamma^\alpha \left(1 - \sup_{\mu, \nu} \lambda_\varepsilon(\mu, \nu) \right) + \sup_{\mu, \nu} \lambda_\varepsilon(\mu, \nu),$$

where

$$\lambda_\varepsilon(\mu, \nu) = \frac{\varepsilon \text{KL}(\pi^* \mid \mu \otimes \nu)}{\bar{W}_{c, \varepsilon}(\mu, \nu)} \in (0, 1)$$

and π^* is the optimal regularized coupling. This implies

$$\bar{W}_{c, \varepsilon}^\infty(\mathcal{T}^\pi \eta, \mathcal{T}^\pi \nu) \leq \bar{\Delta}_\varepsilon(\gamma, \alpha) \bar{W}_{c, \varepsilon}^\infty(\eta, \nu).$$

The rate lies between the unregularized Wasserstein rate also discussed in the book and the regularization dominated rate:

$$\bar{\Delta}_\varepsilon(\gamma, \alpha) \in (\gamma^\alpha, 1).$$

The particle learning rule has the same Bellman target construction as MMD-DQN. With current empirical measure

$$\eta_\omega(x, a) = \frac{1}{N} \sum_{i=1}^N \delta_{z_i(x, a; \omega)}$$

and target empirical measure

$$\hat{\eta}_{\bar{\omega}}(x, a, r, x') = \frac{1}{N} \sum_{j=1}^N \delta_{r + \gamma z_j(x', a^*(x'); \bar{\omega})},$$

the loss is

$$L_{\text{RW}}(\omega) = \overline{W}_{c,\varepsilon}(\eta_\omega(x, a), \hat{\eta}_\omega(x, a, r, x')) .$$

The Sinkhorn algorithm, which allows one to calculate the regularized Wasserstein distance between two measure with finite support, gives a differentiable approximation to this, which makes it usable as a loss. The main downside of having to use this sinkhorn algorithm, is that it is computationally more expensive than standard losses and even the MMD one, but as we will it will also yield stronger result.

Algorithm	Distribution divergence d_p	Representation Z_θ	Contraction parameter of $\mathcal{T}^\pi$
C51	Cramér distance	Categorical distribution	$\sqrt{\gamma}$
QR-DQN	Wasserstein distance	Quantiles	γ
MMD-DQN	MMD	Samples	$\gamma^{\alpha/2}$ for k_α
SinkhornDRL (ours)	Sinkhorn divergence	Samples	γ^α as $\varepsilon \rightarrow 0$; $\overline{\Delta}_\varepsilon(\gamma, \alpha)$ for $0 < \varepsilon < \infty$; γ^α as $\varepsilon \rightarrow \infty$

Table 1: Overview of the Distributional RL algorithms we discussed here. (Table is taken from [Sun+24].)

Theorem 8: Wasserstein–MMD interpolation [Sun+24, Theorem 1]

The regularization parameter interpolates between two geometries:

$$\varepsilon \rightarrow 0 \quad \Rightarrow \quad \overline{W}_{c,\varepsilon}(\mu, \nu) \rightarrow 2W_c^\alpha(\mu, \nu),$$

whereas

$$\varepsilon \rightarrow +\infty \quad \Rightarrow \quad \overline{W}_{c,\varepsilon}(\mu, \nu) \rightarrow \text{MMD}_{k_\alpha}^2(\mu, \nu), \quad k_\alpha = -c.$$

Thus the small-regularization endpoint recovers transport geometry, while the large-regularization endpoint recovers the MMD induced by the negative cost kernel.

Proof. The proof can be found in [GPC18], the main argument is treating it as a constraint problem and looking at the dual. □

This, in a sense also shows that the sinkhornDLR paper is a strict generalization of the MMD paper. At least when taking the kernel to be $-c$.

Finally, we also report some of the results form [Sun+24] on the Atari benchmark comparing the Sinkhorn method with both MMD and the classical C51 and quantile methods.




Figure 4: Atari performance for all methods covered. The three panels show the mean, median, and interquartile mean of human-normalized scores over training. Results taken from [BDR23]

References

- [BDR23] Marc G. Bellemare, Will Dabney, and Mark Rowland. *Distributional Reinforcement Learning*. MIT Press, 2023.
- [GPC18] Aude Genevay, Gabriel Peyré, and Marco Cuturi. “Learning Generative Models with Sinkhorn Divergences”. In: *Proceedings of the Twenty-First International Conference on Artificial Intelligence and Statistics*. Vol. 84. Proceedings of Machine Learning Research. PMLR, 2018, pp. 1608–1617.
- [NTGV21] Thanh Nguyen-Tang, Sunil Gupta, and Svetha Venkatesh. “Distributional Reinforcement Learning via Moment Matching”. In: *Proceedings of the AAAI Conference on Artificial Intelligence*. Vol. 35. 10. 2021, pp. 9144–9152. DOI: 10.1609/aaai.v35i10.17104.
- [Sej+13] Dino Sejdinovic et al. “Equivalence of distance-based and RKHS-based statistics in hypothesis testing”. In: *The annals of statistics* (2013), pp. 2263–2291.
- [Sun+24] Ke Sun et al. “Distributional Reinforcement Learning with Regularized Wasserstein Loss”. In: *Advances in Neural Information Processing Systems*. Vol. 37. 2024.